

# Spam Emails

- Mayank Srivastava



## About Dataset

### Overview:

This dataset contains a collection of emails, categorized into two classes: "Spam" and "Non-Spam" (often referred to as "Ham"). These emails have been carefully curated and labeled to aid in the development of spam email detection models. Whether you are interested in email filtering, natural language processing, or machine learning, this dataset can serve as a valuable resource for training and evaluation.

### Context:

Spam emails continue to be a significant issue, with malicious actors attempting to deceive users with unsolicited, fraudulent, or harmful messages. This dataset is designed to facilitate research, development, and testing of algorithms and models aimed at accurately identifying and filtering spam emails, helping protect users from various threats.

### Content:

The dataset includes the following features: Message: The content of the email, including the subject line and message body, Category: Categorizes each email as either "Spam" or "Ham" (Non-Spam).

### Potential Use Cases:

1. Email Filtering: Develop and evaluate email filtering systems that automatically classify incoming emails as spam or non-spam.
2. Natural Language Processing (NLP): Use the email text for text classification, topic modeling, and sentiment analysis.
3. Machine Learning: Create machine learning models for spam detection, potentially employing various algorithms and techniques.
4. Feature Engineering: Explore email content features that contribute to spam classification accuracy.
5. Data Analysis: Investigate patterns and trends in spam email content and characteristics.

### License:

Please note that this dataset is for research and analysis purposes only and may be subject to copyright and data use restrictions. Ensure compliance with relevant policies when using this data.

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
        5 df= pd.read_csv('spam.csv')
        6 df.head()
```

Out[1]:

|   | Category | Message   |
|---|----------|---|
| 0 | ham      | Go until jurong point, crazy.. Available only ... |
| 1 | ham      | Ok lar... Joking wif u oni...                     |
| 2 | spam     | Free entry in 2 a wkly comp to win FA Cup fina... |
| 3 | ham      | U dun say so early hor... U c already then say... |
| 4 | ham      | Nah I don't think he goes to usf, he lives aro... |

```
In [2]: 1 #checking the shape
        2 df.shape
```

Out[2]: (5572, 2)

```
In [3]: 1 # checking the info()
        2 df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Category    5572 non-null   object
1   Message     5572 non-null   object
dtypes: object(2)
memory usage: 87.2+ KB
```

### Observations:

1. shape = 5572 rows x 2 columns
2. target column= Category
3. Null values = NA

```
In [4]: 1 data=df.Category.value_counts()
2 plt.pie(data, labels=data.index, autopct='%1.2f%%', explode=[0,0.1])
3 plt.title('Percentage Spam vs Ham mails in the Dataset')
4 plt.show()
```

Percentage Spam vs Ham mails in the Dataset



```
In [5]: 1 # checking Category Distribution
2 df.Category.value_counts()
```

```
Out[5]: Category
ham      4825
spam      747
Name: count, dtype: int64
```

## Standardization

```
In [6]: 1 import spacy
2 nlp=spacy.load('en_core_web_sm')
```

```
In [7]: 1 import re
2 from nltk import stem
3 def standardize(text):
4     #converted to lowercase
5     text=text.lower()
6
7     #removed punctuations
8     doc=re.findall(r'\w+', text)
9     text=' '.join(doc)
10
11     #Lemmatize
12     doc =nlp(text)
13     text = ' '.join([token.lemma_ for token in doc])
14
15     #stem
16     ps=stem.PorterStemmer()
17     text=' '.join([ps.stem(token) for token in text.split(' ')])
18
19     return text
```

```
In [8]: 1 df.Message =df.Message.apply(standardize)
```

```
In [9]: 1 df.Message
```

```
Out[9]: 0      go until jurong point crazi avail onli in bugi...
1      ok lar joke wif u oni
2      free entri in 2 a wkli comp to win fa cup fina...
3      u dun say so earli hor u c alreadi then say
4      nah i don t think he go to usf he live around ...
...
5567     thi be the 2nd time we have tri 2 contact u u ...
5568     will ü b go to esplanad fr home
5569     piti be in mood for that so ani other suggest
5570     the guy do some bitch but i act like i d be in...
5571     rofl it true to it name
Name: Message, Length: 5572, dtype: object
```

```
In [10]: 1 # Encoding the target class
2 # spam =1 and ham =0
3 df.Category =df.Category.replace(['ham','spam'], [0,1])
```

## train\_test\_split

```
In [11]: 1 from sklearn.model_selection import train_test_split
2 x=df.Message
3 y=df.Category
4 xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=0.3,random_state=40, stratify = y)
```

## Vectorization

```
In [12]: 1 from sklearn.feature_extraction.text import TfidfVectorizer
```

```
In [13]: 1 tv=TfidfVectorizer(stop_words='english')
2 xtrain_tv=tv.fit_transform(xtrain)
3 xtest_tv=tv.transform(xtest)
```

## Training ML models

```
In [14]: 1 from sklearn.linear_model import LogisticRegression
2 from sklearn.tree import DecisionTreeClassifier
3 from sklearn.svm import SVC
4 from sklearn.neighbors import KNeighborsClassifier
5 from sklearn.ensemble import AdaBoostClassifier,RandomForestClassifier
6 from sklearn.naive_bayes import GaussianNB
```

```
In [15]: 1 logr= LogisticRegression()
2 knn=KNeighborsClassifier()
3 dt=DecisionTreeClassifier()
4 rf=RandomForestClassifier()
5 svc=SVC()
6 ad=AdaBoostClassifier()
7 nb=GaussianNB()
```

```
In [16]: 1 models = {'LOGR': logr, "DT": dt, 'RF': rf, 'AD': ad, 'SVC': svc, 'KNN':knn, 'NB':nb}
```

```
In [17]: 1 #metrics
2 from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score, confusion_matrix
```

```
In [*]: 1 for i in models:
2     model=models[i]
3     model.fit(xtrain_tv.toarray(), ytrain)
4     print(model)
5
6     #predicting training data
7     predict_xtrain = model.predict(xtrain_tv.toarray())
8     print('training accuracy_score: ', accuracy_score(ytrain, predict_xtrain))
9     print('training precision_score: ', precision_score(ytrain, predict_xtrain))
10    print('training recall_score: ', recall_score(ytrain, predict_xtrain))
11    print('training f1_score: ', f1_score(ytrain, predict_xtrain))
12    print('training confusion_matrix: \n', confusion_matrix(ytrain, predict_xtrain))
13
14    #predicting testing data
15    predict_xtest = model.predict(xtest_tv.toarray())
16    print('testing accuracy_score: ', accuracy_score(ytest, predict_xtest))
17    print('testing precision_score: ', precision_score(ytest, predict_xtest))
18    print('testing recall_score: ', recall_score(ytest, predict_xtest))
19    print('testing f1_score: ', f1_score(ytest, predict_xtest))
20    print('testing confusion_matrix: \n', confusion_matrix(ytest, predict_xtest))
21    print('\n\n')
```

LogisticRegression()  
training accuracy\_score: 0.9664102564102565  
training precision\_score: 0.9949494949494949  
training recall\_score: 0.7533460803059273  
training f1\_score: 0.8574537540805223  
training confusion\_matrix:  
[[3375 2]  
 [129 394]]  
testing accuracy\_score: 0.9593301435406698  
testing precision\_score: 0.9588235294117647  
testing recall\_score: 0.7276785714285714  
testing f1\_score: 0.8274111675126904  
testing confusion\_matrix:  
[[1441 7]  
 [ 61 163]]

DecisionTreeClassifier()  
training accuracy\_score: 1.0  
training precision\_score: 1.0  
training recall\_score: 1.0  
training f1\_score: 1.0  
training confusion\_matrix:  
[[3377 0]  
 [ 0 523]]  
testing accuracy\_score: 0.9665071770334929  
testing precision\_score: 0.8716814159292036  
testing recall\_score: 0.8794642857142857  
testing f1\_score: 0.8755555555555555  
testing confusion\_matrix:  
[[1419 29]  
 [ 27 197]]

RandomForestClassifier()  
training accuracy\_score: 0.9997435897435898  
training precision\_score: 1.0  
training recall\_score: 0.9980879541108987  
training f1\_score: 0.999043062200957  
training confusion\_matrix:  
[[3377 0]  
 [ 1 522]]  
testing accuracy\_score: 0.9796650717703349  
testing precision\_score: 0.9656862745098039  
testing recall\_score: 0.8794642857142857  
testing f1\_score: 0.9205607476635514  
testing confusion\_matrix:  
[[1441 7]  
 [ 27 197]]

AdaBoostClassifier()  
training accuracy\_score: 0.9805128205128205  
training precision\_score: 0.976545842217484  
training recall\_score: 0.875717017208413  
training f1\_score: 0.9233870967741935  
training confusion\_matrix:  
[[3366 11]  
 [ 65 458]]  
testing accuracy\_score: 0.9742822966507177  
testing precision\_score: 0.9547738693467337  
testing recall\_score: 0.8482142857142857  
testing f1\_score: 0.8983451536643027  
testing confusion\_matrix:  
[[1439 9]  
 [ 34 190]]

|                          | Prediction Accuracy (%) |         | Testing Sample Performance |        |          |
|--------------------------|-------------------------|---------|----------------------------|--------|----------|
| Model                    | Training                | Testing | Precision                  | Recall | F1 score |
| LogisticRegression()     | 96.64                   | 95.93   | 95.88                      | 72.76  | 82.74    |
| DecisionTreeClassifier() | 100                     | 96.77   | 87.94                      | 87.94  | 87.94    |
| RandomForestClassifier() | 100                     | 97.66   | 96.48                      | 85.71  | 90.78    |
| AdaBoostClassifier()     | 98.05                   | 97.3    | 94.97                      | 84.37  | 89.36    |
| SVC()                    | 99.79                   | 98.2    | 97.54                      | 88.83  | 92.99    |
| KNeighborsClassifier()   | 91.46                   | 90.01   | 98.3                       | 25.89  | 40.98    |
| GaussianNB()             | 91.02                   | 86.48   | 49.75                      | 90.62  | 64.24    |